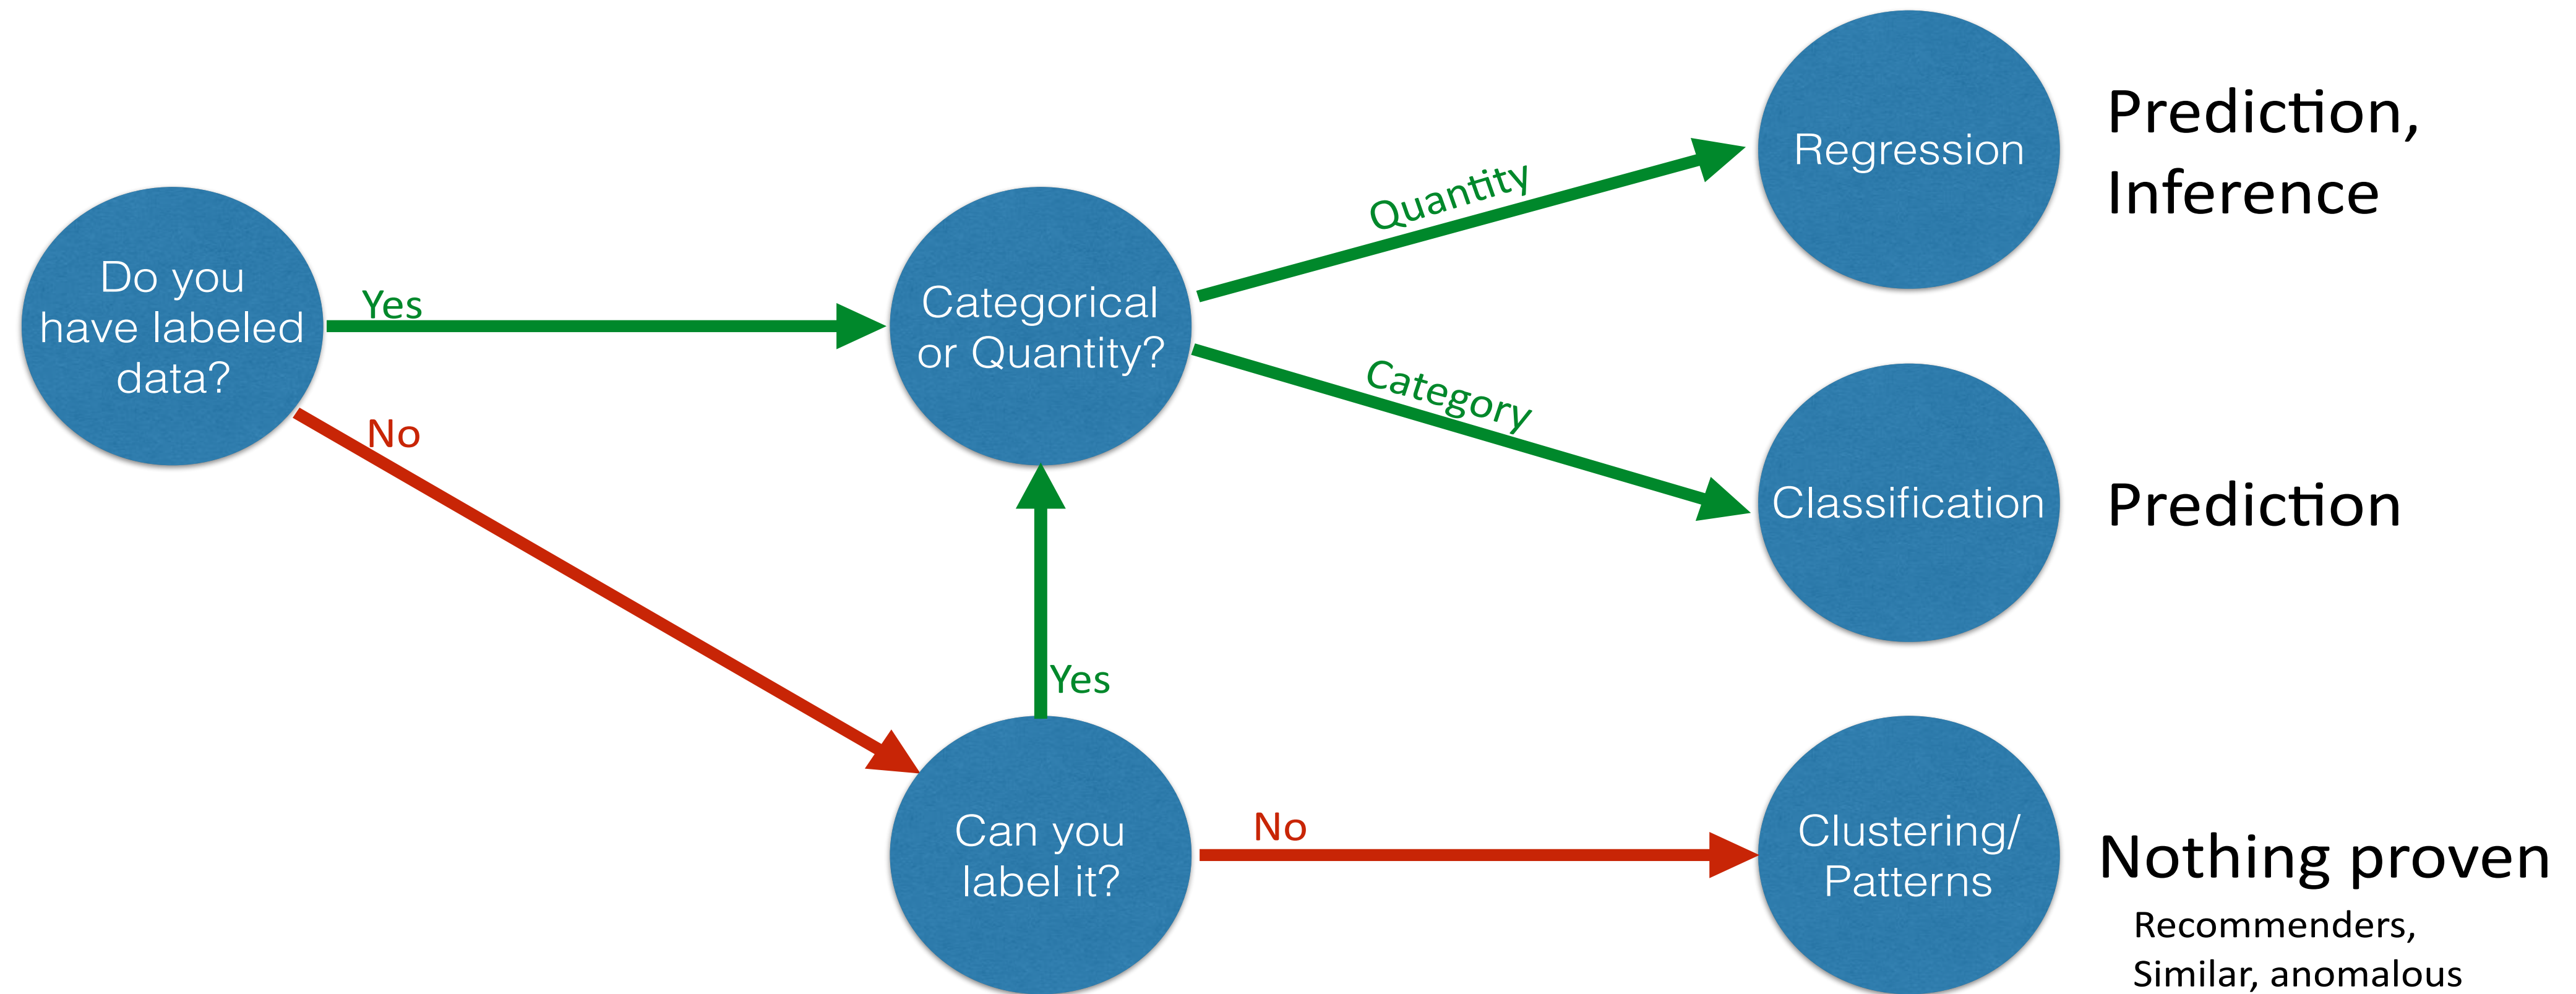


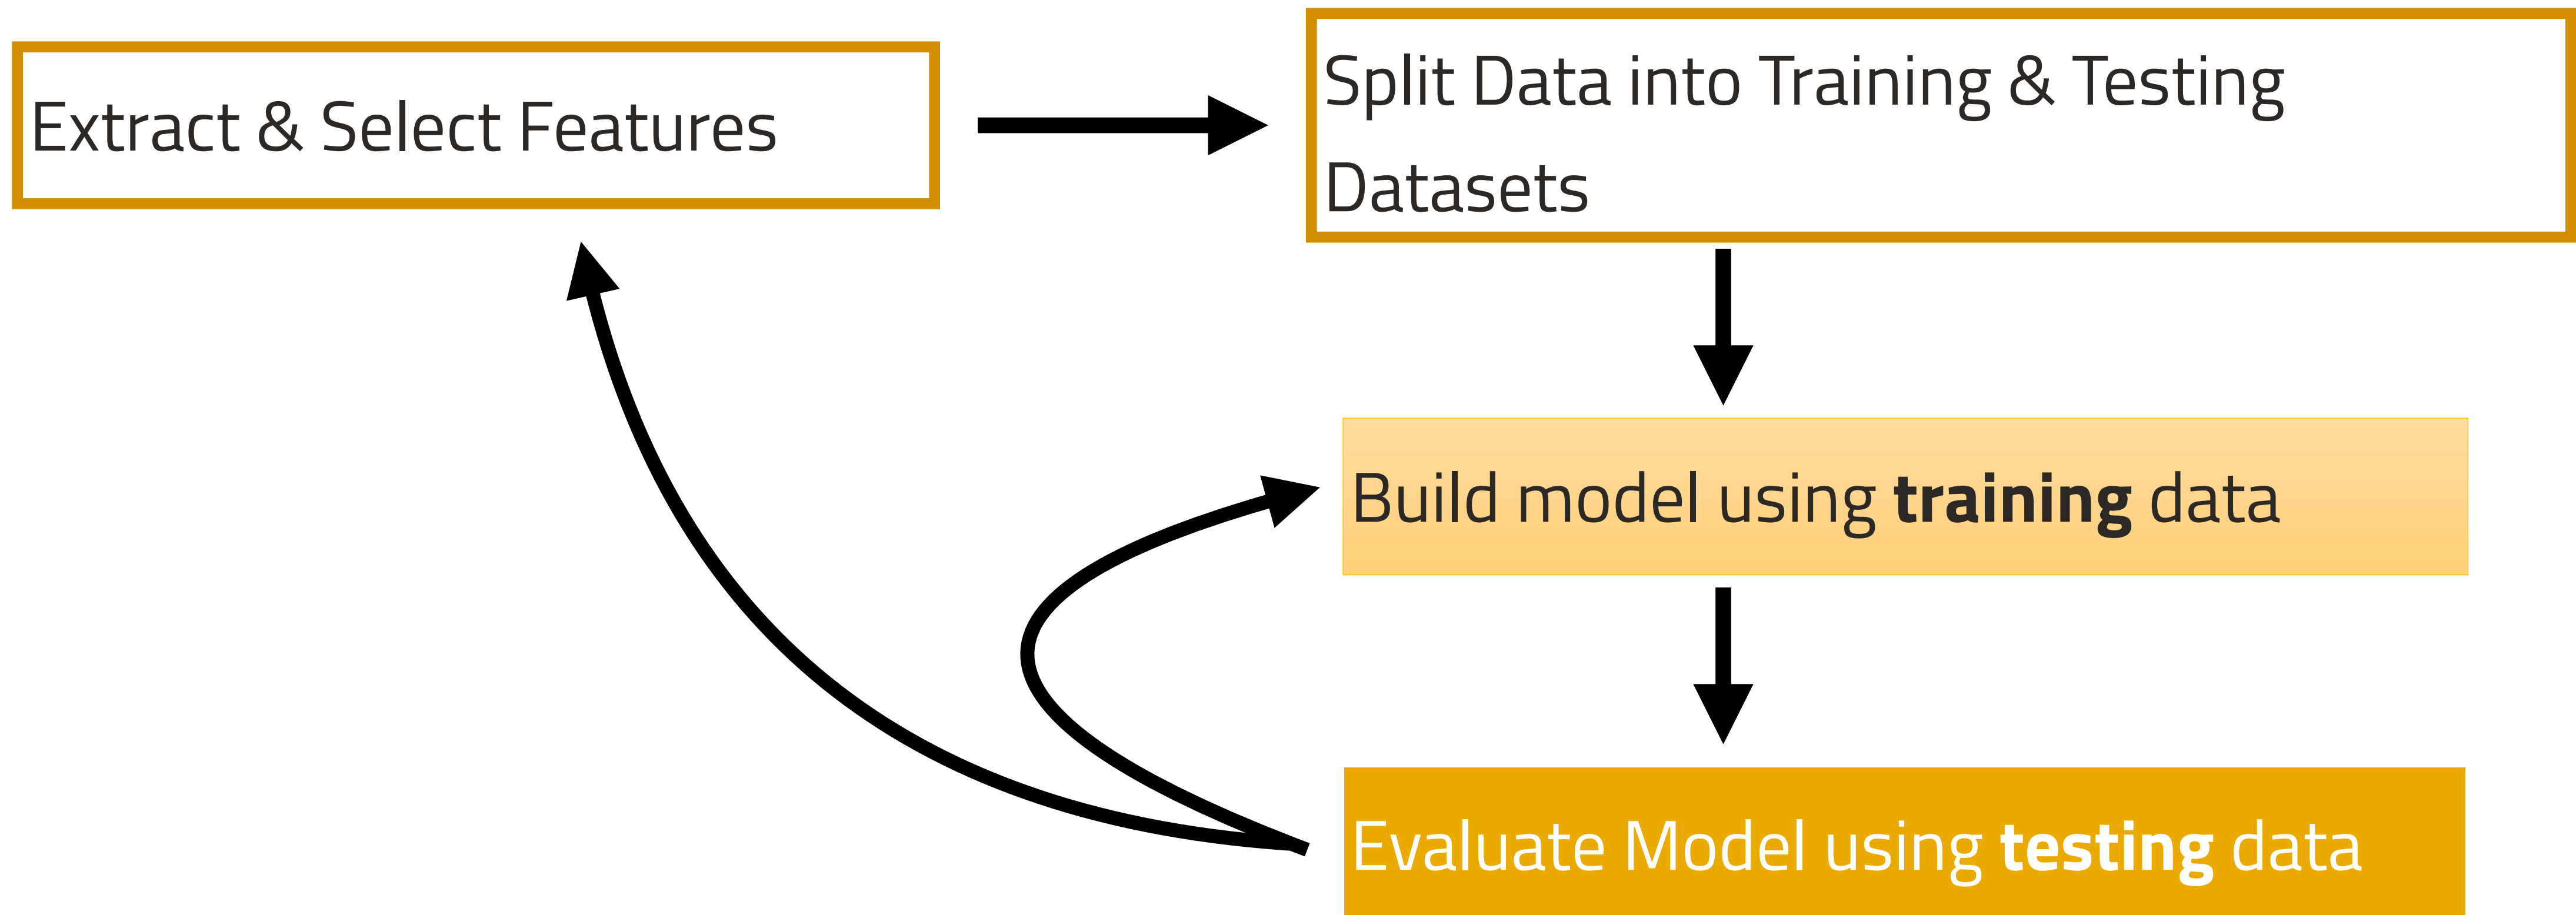
Module 6: Unsupervised Learning: Clustering

Agenda for Today

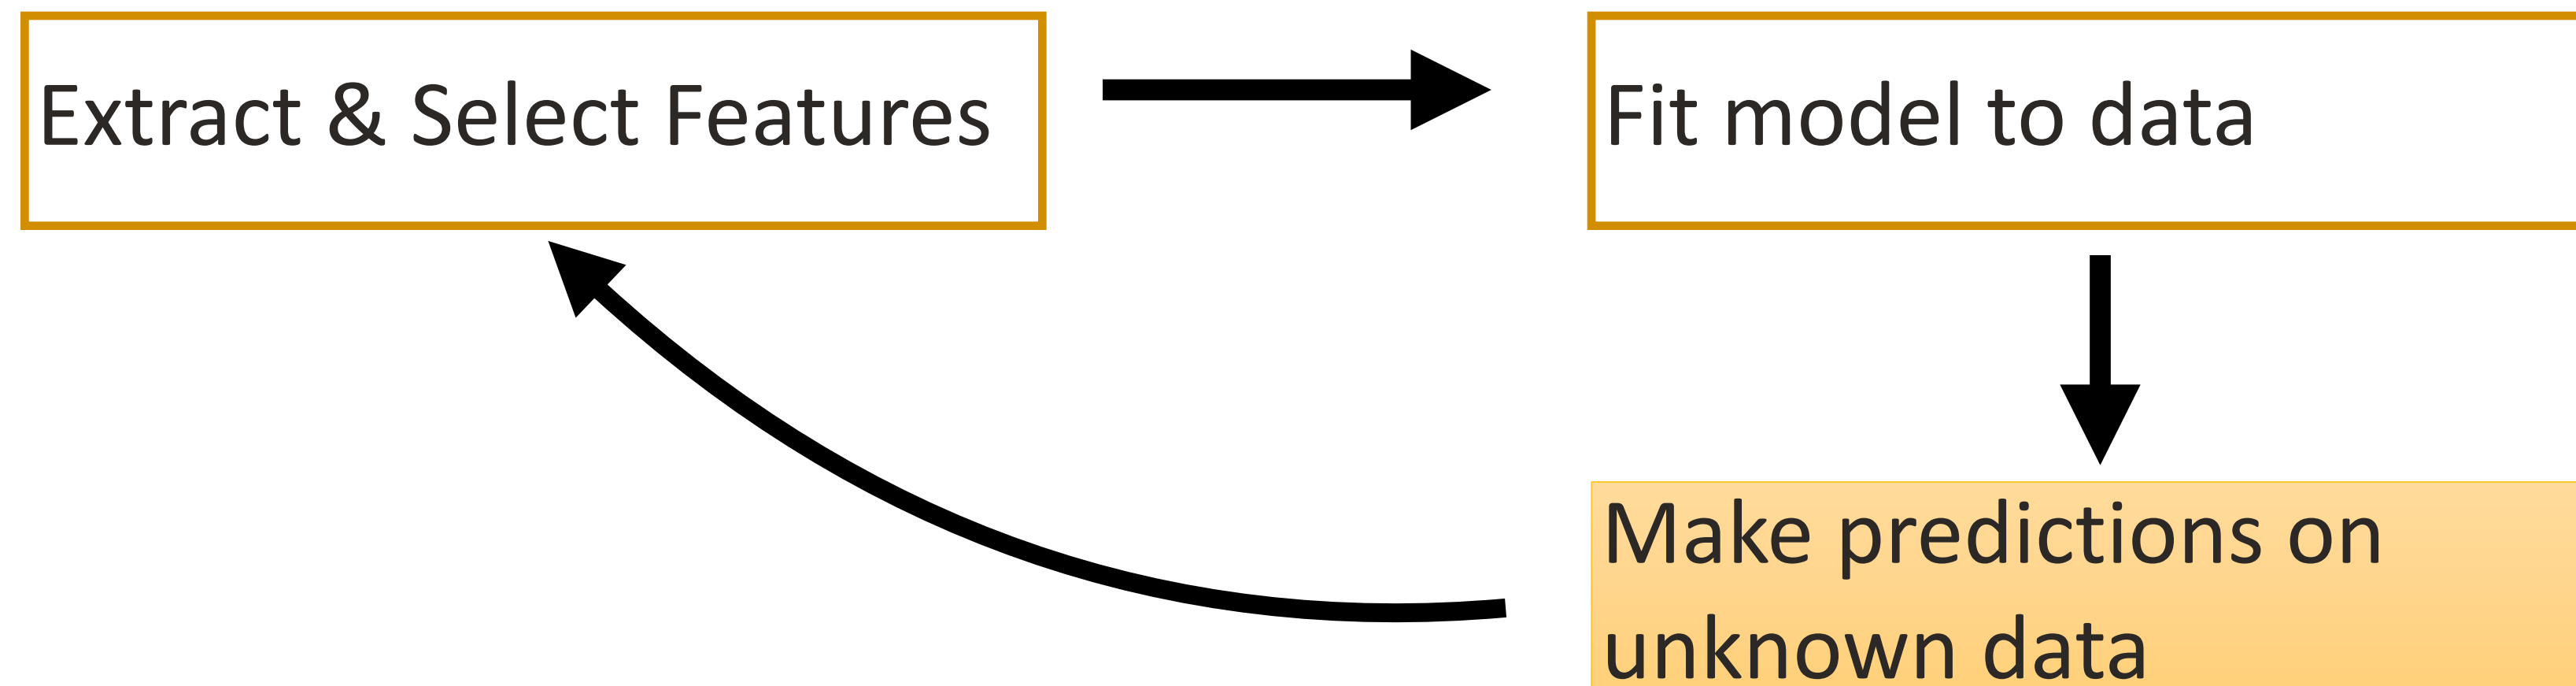
- Measuring Distances
- Math free overview of clustering techniques



Supervised ML Process



Unsupervised ML Process



Unsupervised Clustering Algorithm

1. Select Features
2. Calculate a distance measure
3. Apply a clustering algorithm
4. Validate?

Which Departments are Similar?

	Malware events
Dept1	6
Dept2	1
Dept3	8

Which Departments are Similar?

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



Which Departments are Similar?

	Malware events	Phishing	Open Tickets
Dept1	6	6	3
Dept2	1	2	1
Dept3	8	1	9

Computing Distance

	Malware events
Dept1	6
Dept2	1
Dept3	8

Compare:

Dept1 to Dept2: $|6 - 1| = 5$

Dept2 to Dept3: $|1 - 8| = 7$

Dept1 to Dept3: $|6 - 8| = 2$

Two-Dimensional Distance

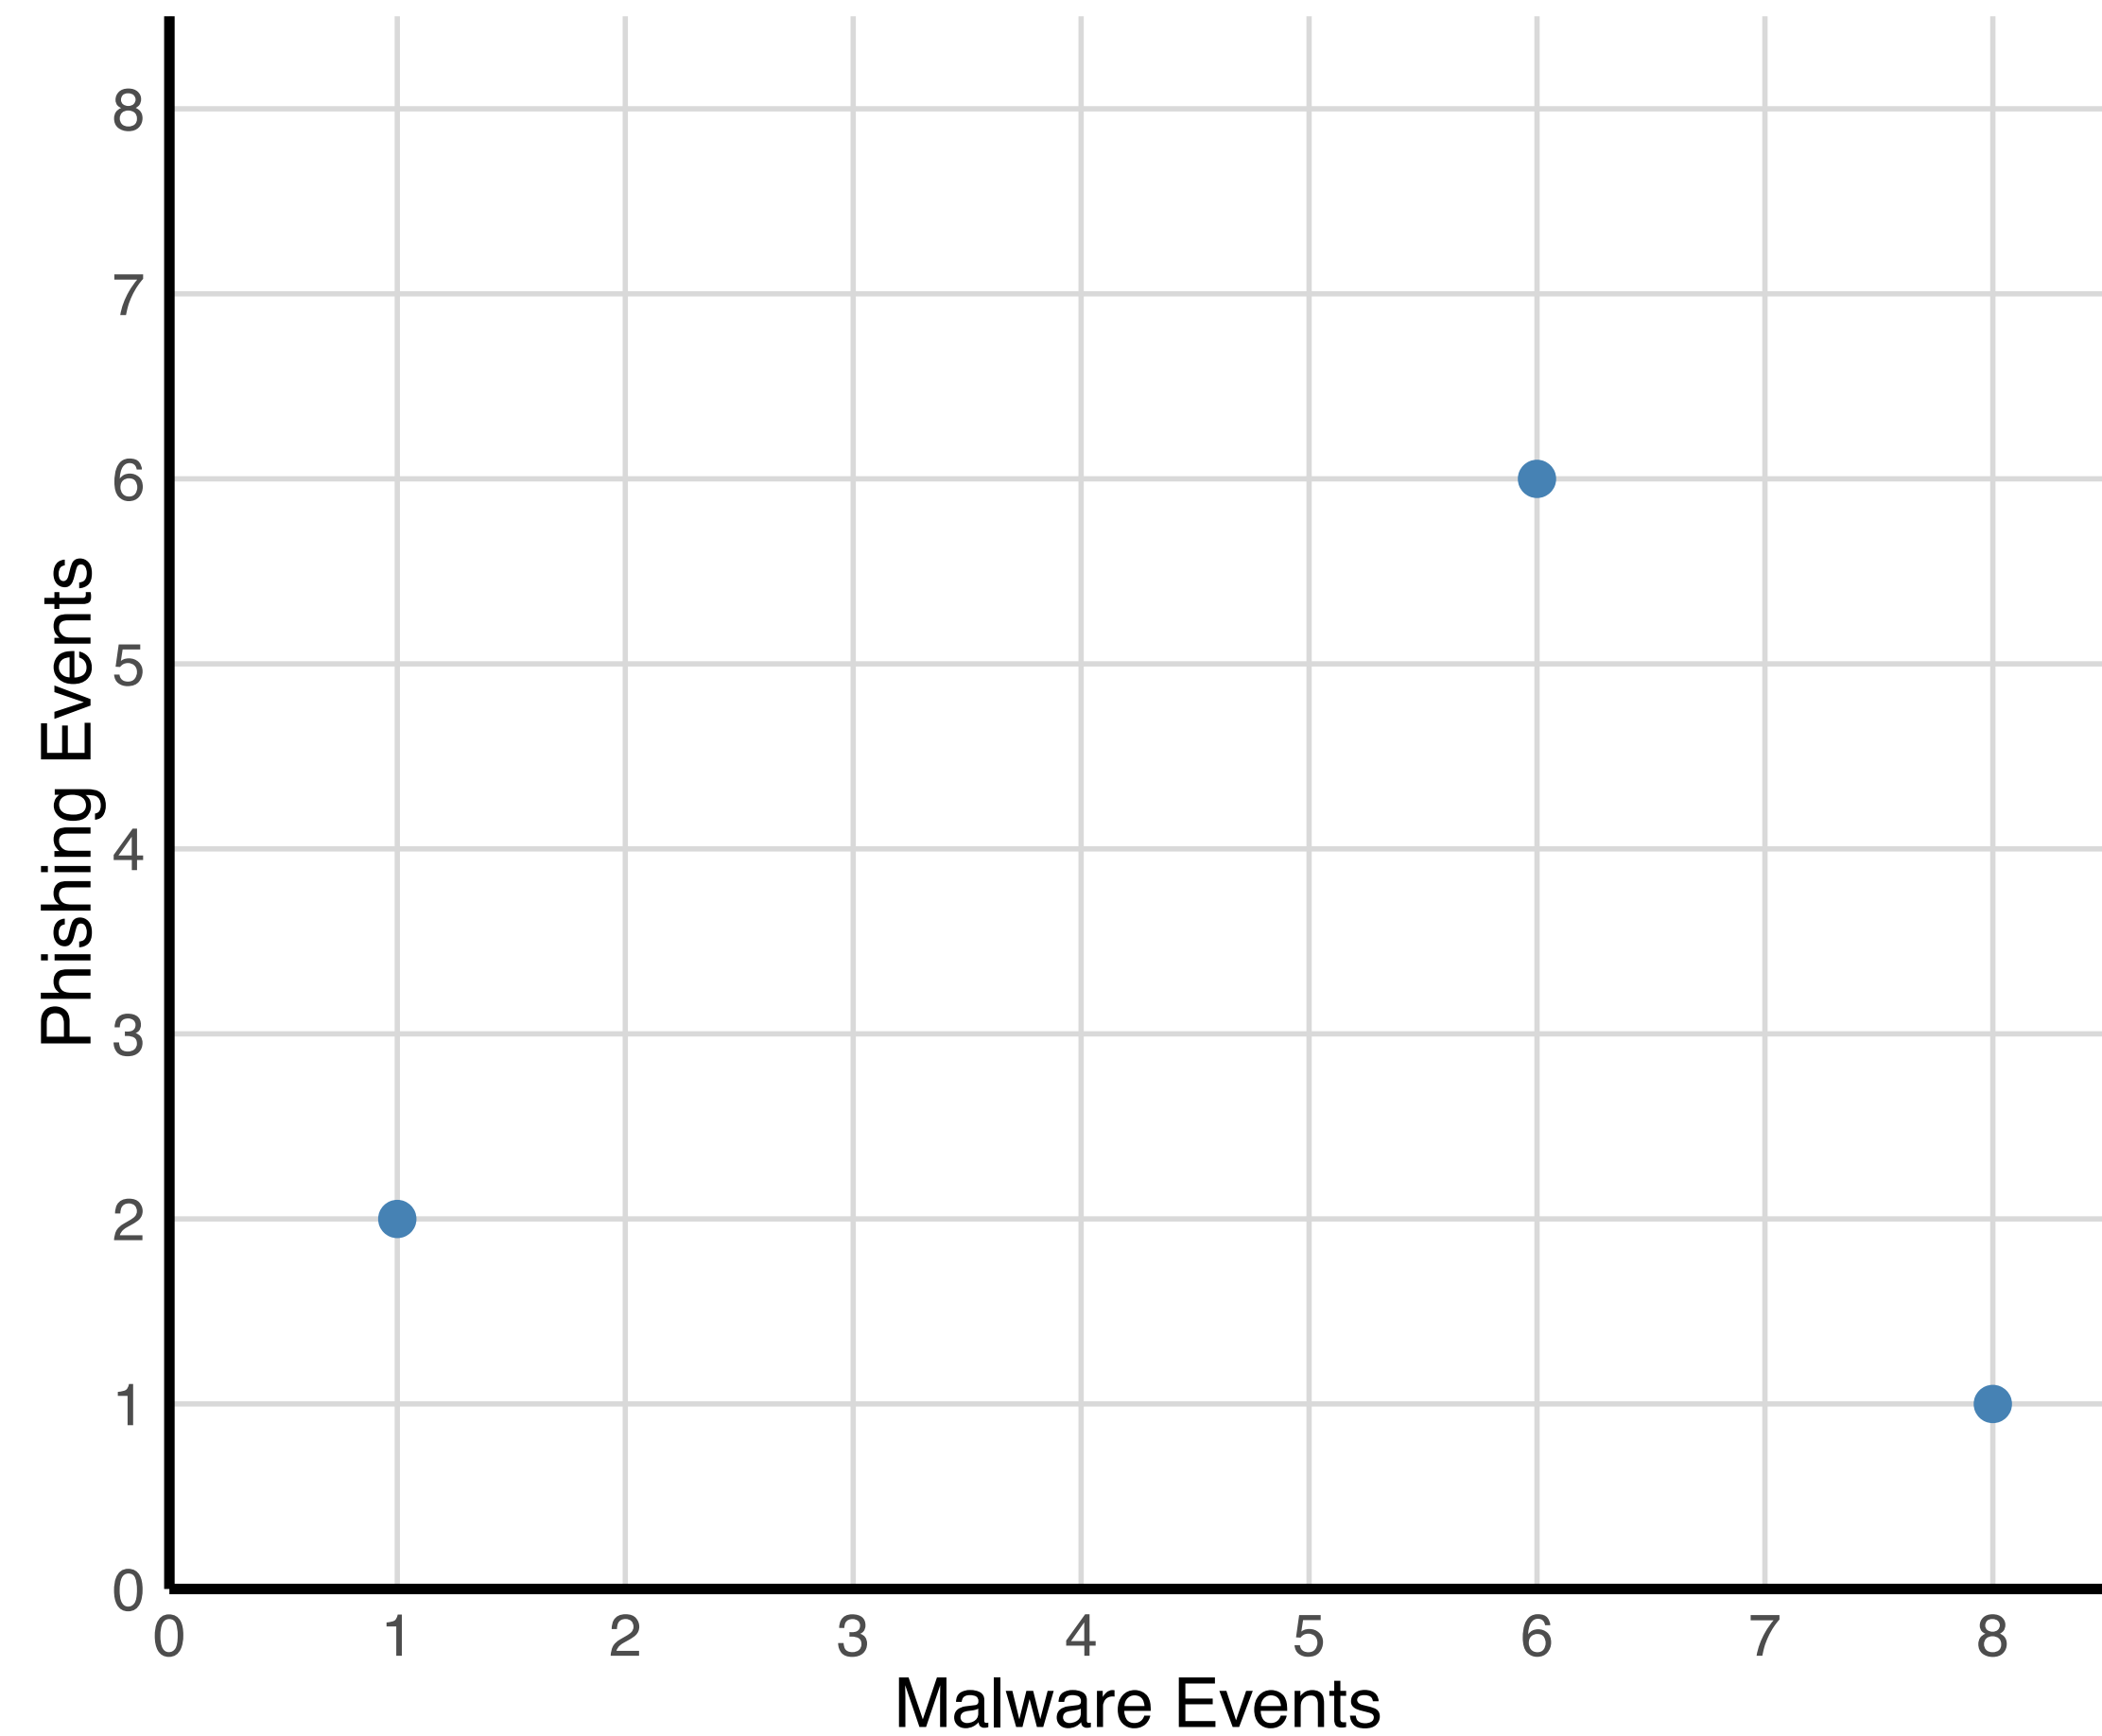
	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1

Multiple Distance methods

- Euclidean
 - Manhattan
 - Maximum
 - Canberra
 - Binary
 - Minkowski
- ... (to name a few)

Two-Dimensional Distance


	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



Two-Dimensional Distance

Euclidean very common and easy to understand

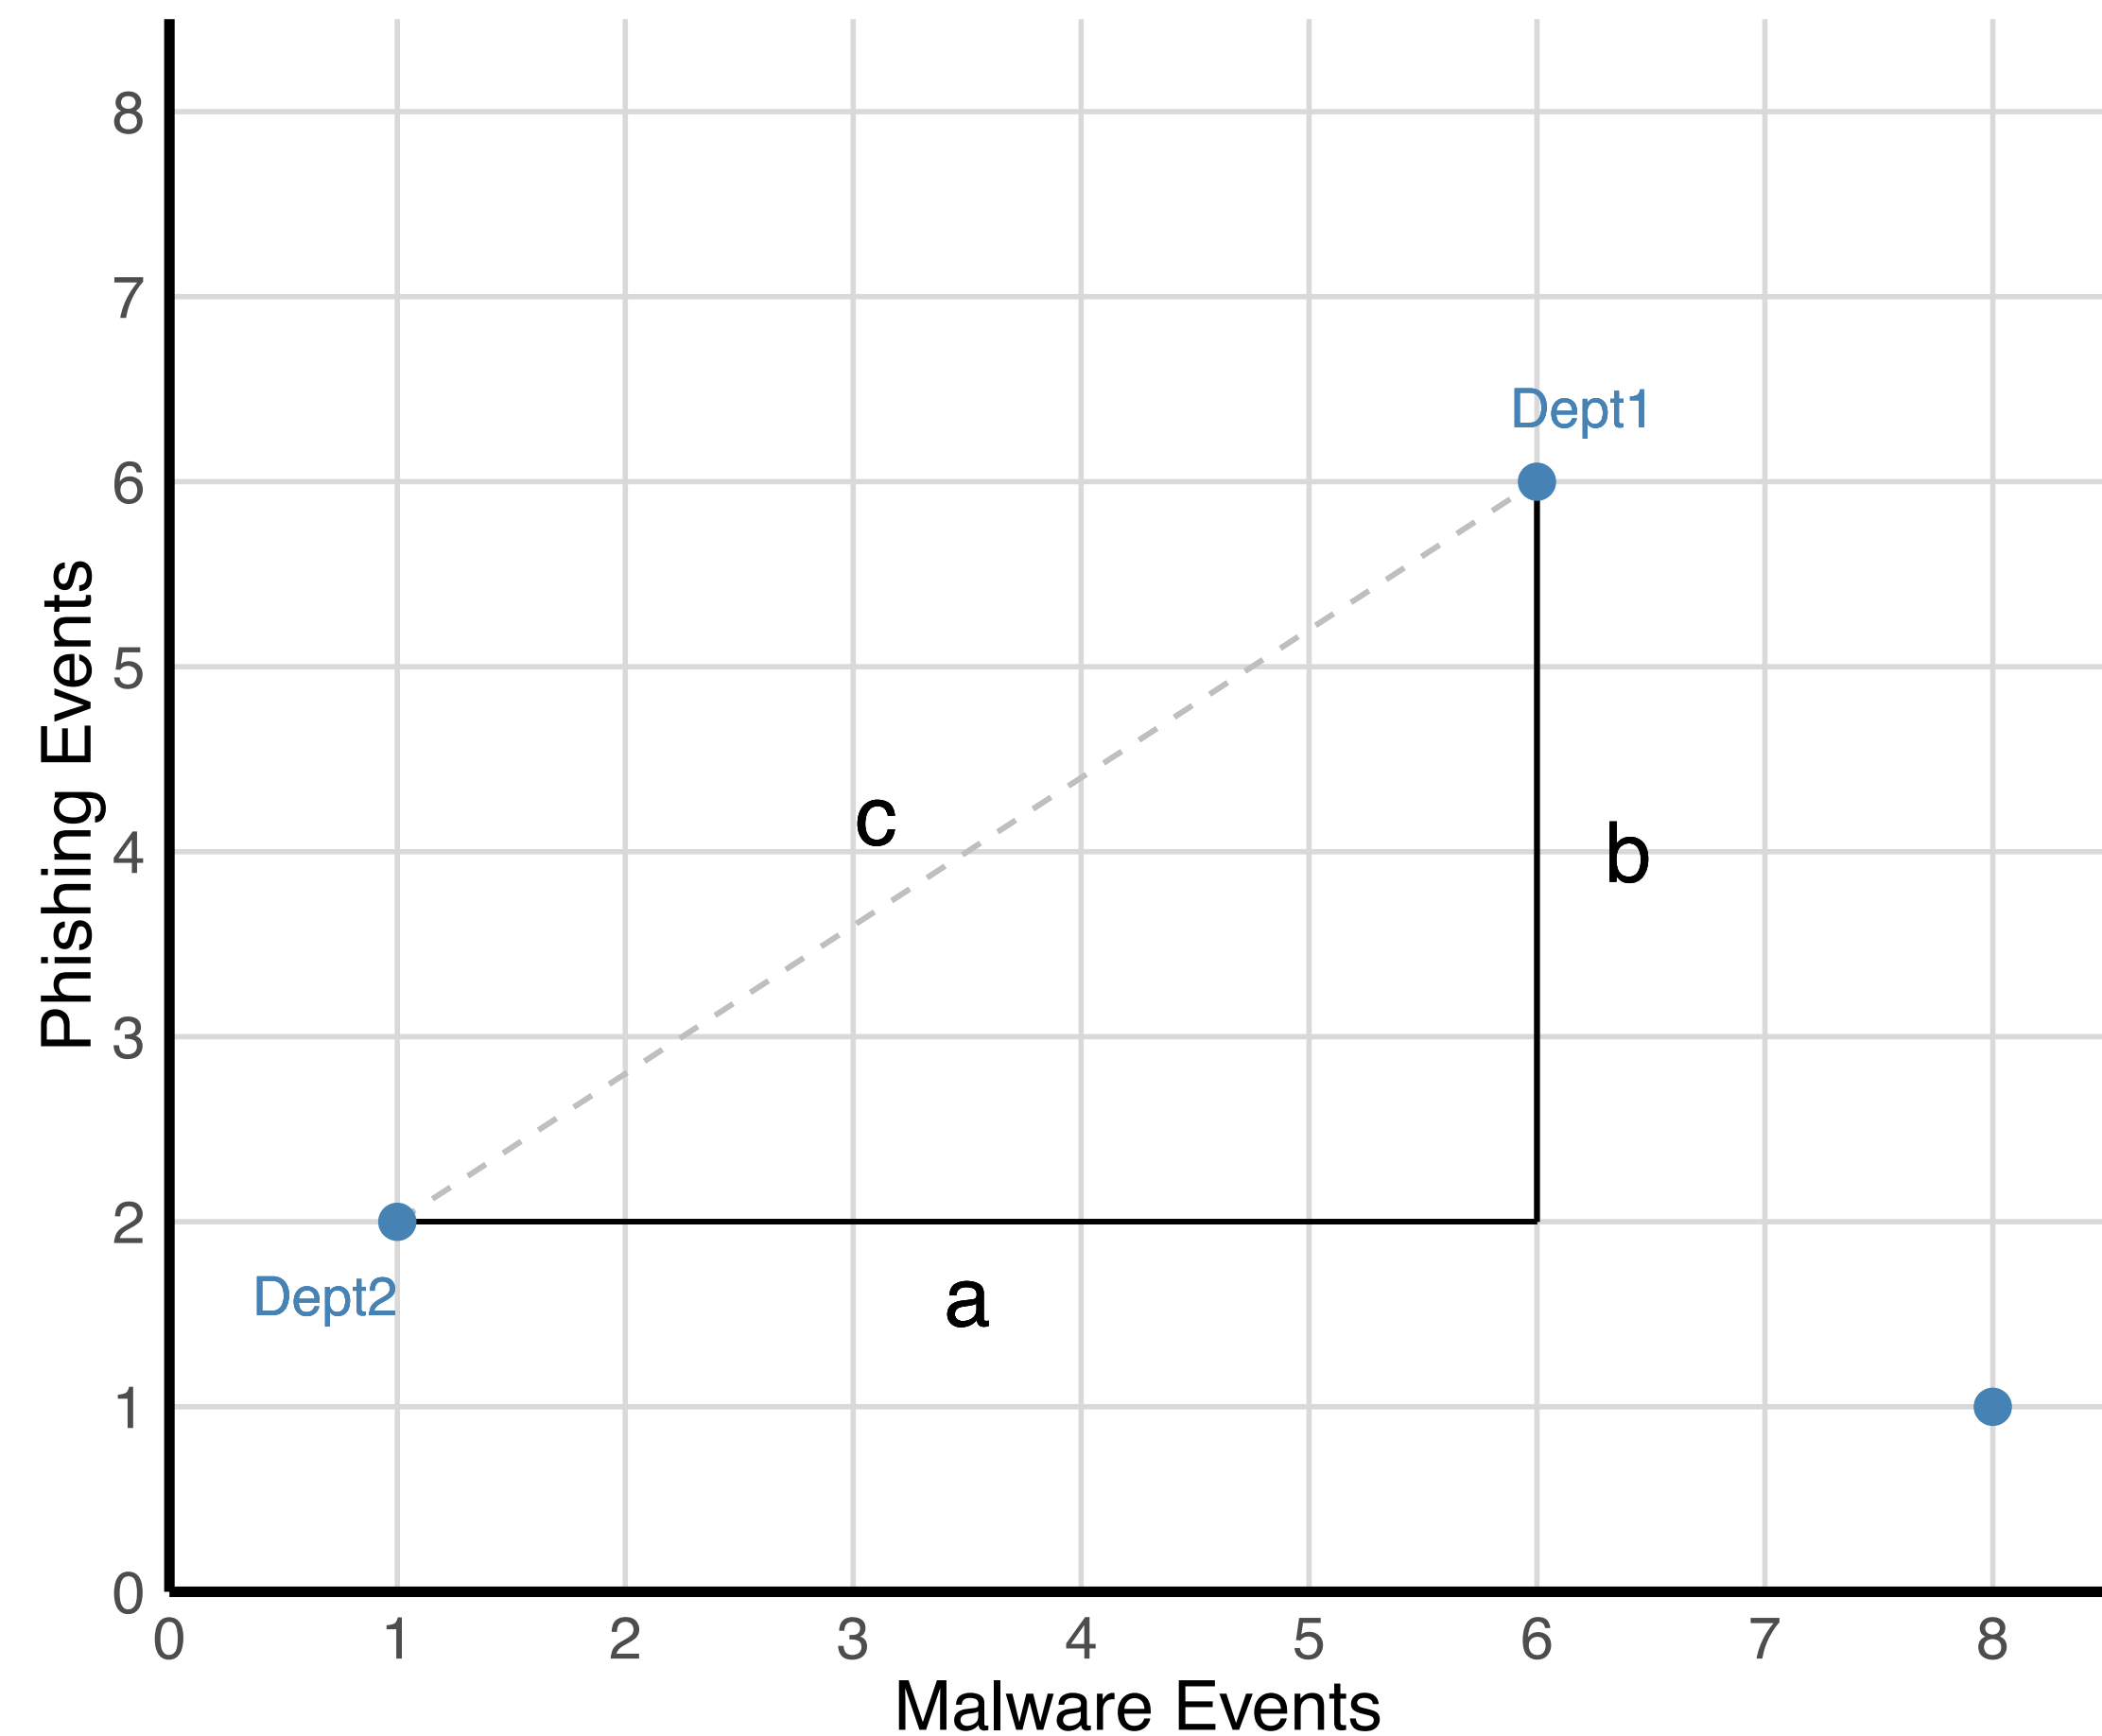
	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



Two-Dimensional Distance

Euclidean very common and easy to understand: $a^2 + b^2 = c^2$


	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



Two-Dimensional Distance

Manhattan also easy to comprehend: $a + b$

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



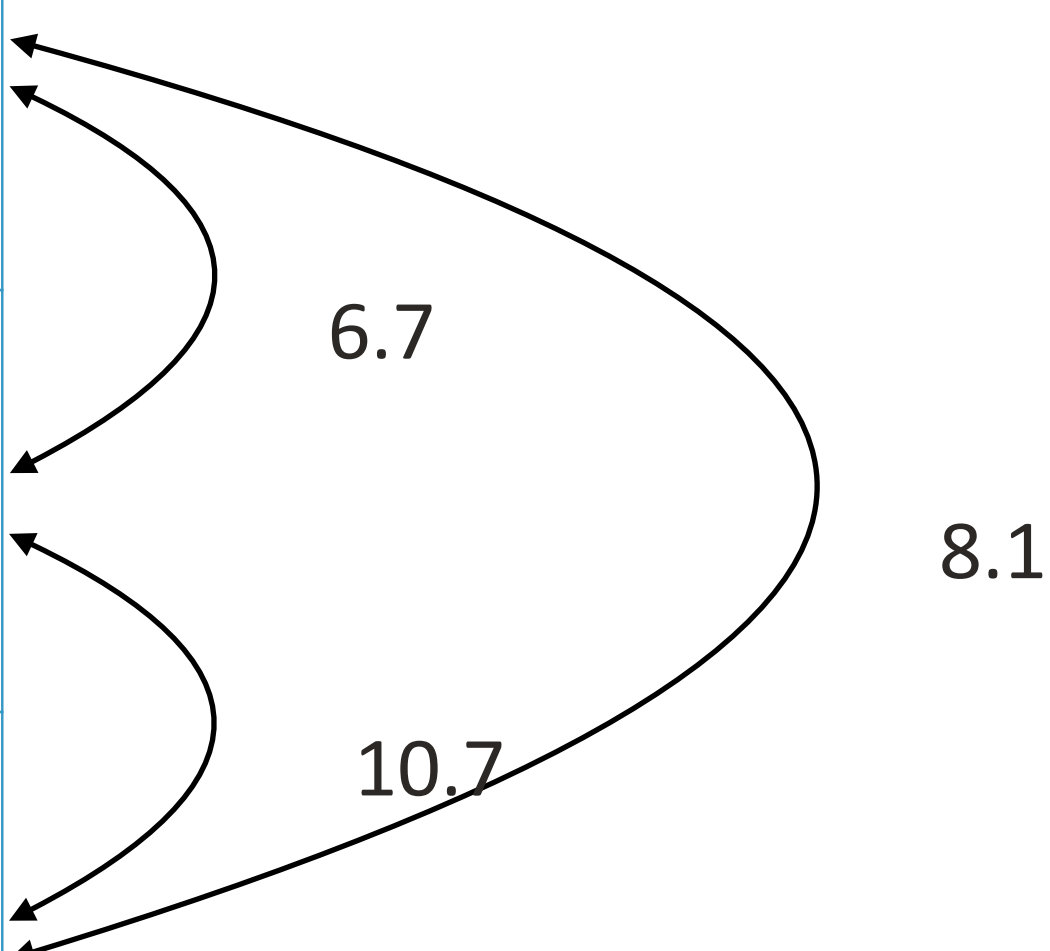
Computing Distance

	Malware events	Phishing	<div>Compare: Dept1 to Dept2: $\sqrt{(6-1)^2 + (6-2)^2} = \mathbf{6.4}$ Dept2 to Dept3: ... = 7.1 Dept1 to Dept3: ... = 5.4</div>
Dept1	6	6	
Dept2	1	2	
Dept3	8	1	

Euclidean Distance calculations


```
def dist(x,y):  
    return np.sqrt(np.sum( (x-y)**2 ))  
  
> mat = np.array([[ 6,6,3 ], [1,2,1], [8,1,9]])  
> dist(mat[0], mat[1])  
6.7082039324993694  
  
> dist(mat[1], mat[2])  
10.677078252031311  
  
> dist(mat[0], mat[2])  
8.0622577482985491
```

Which Departments are Similar?

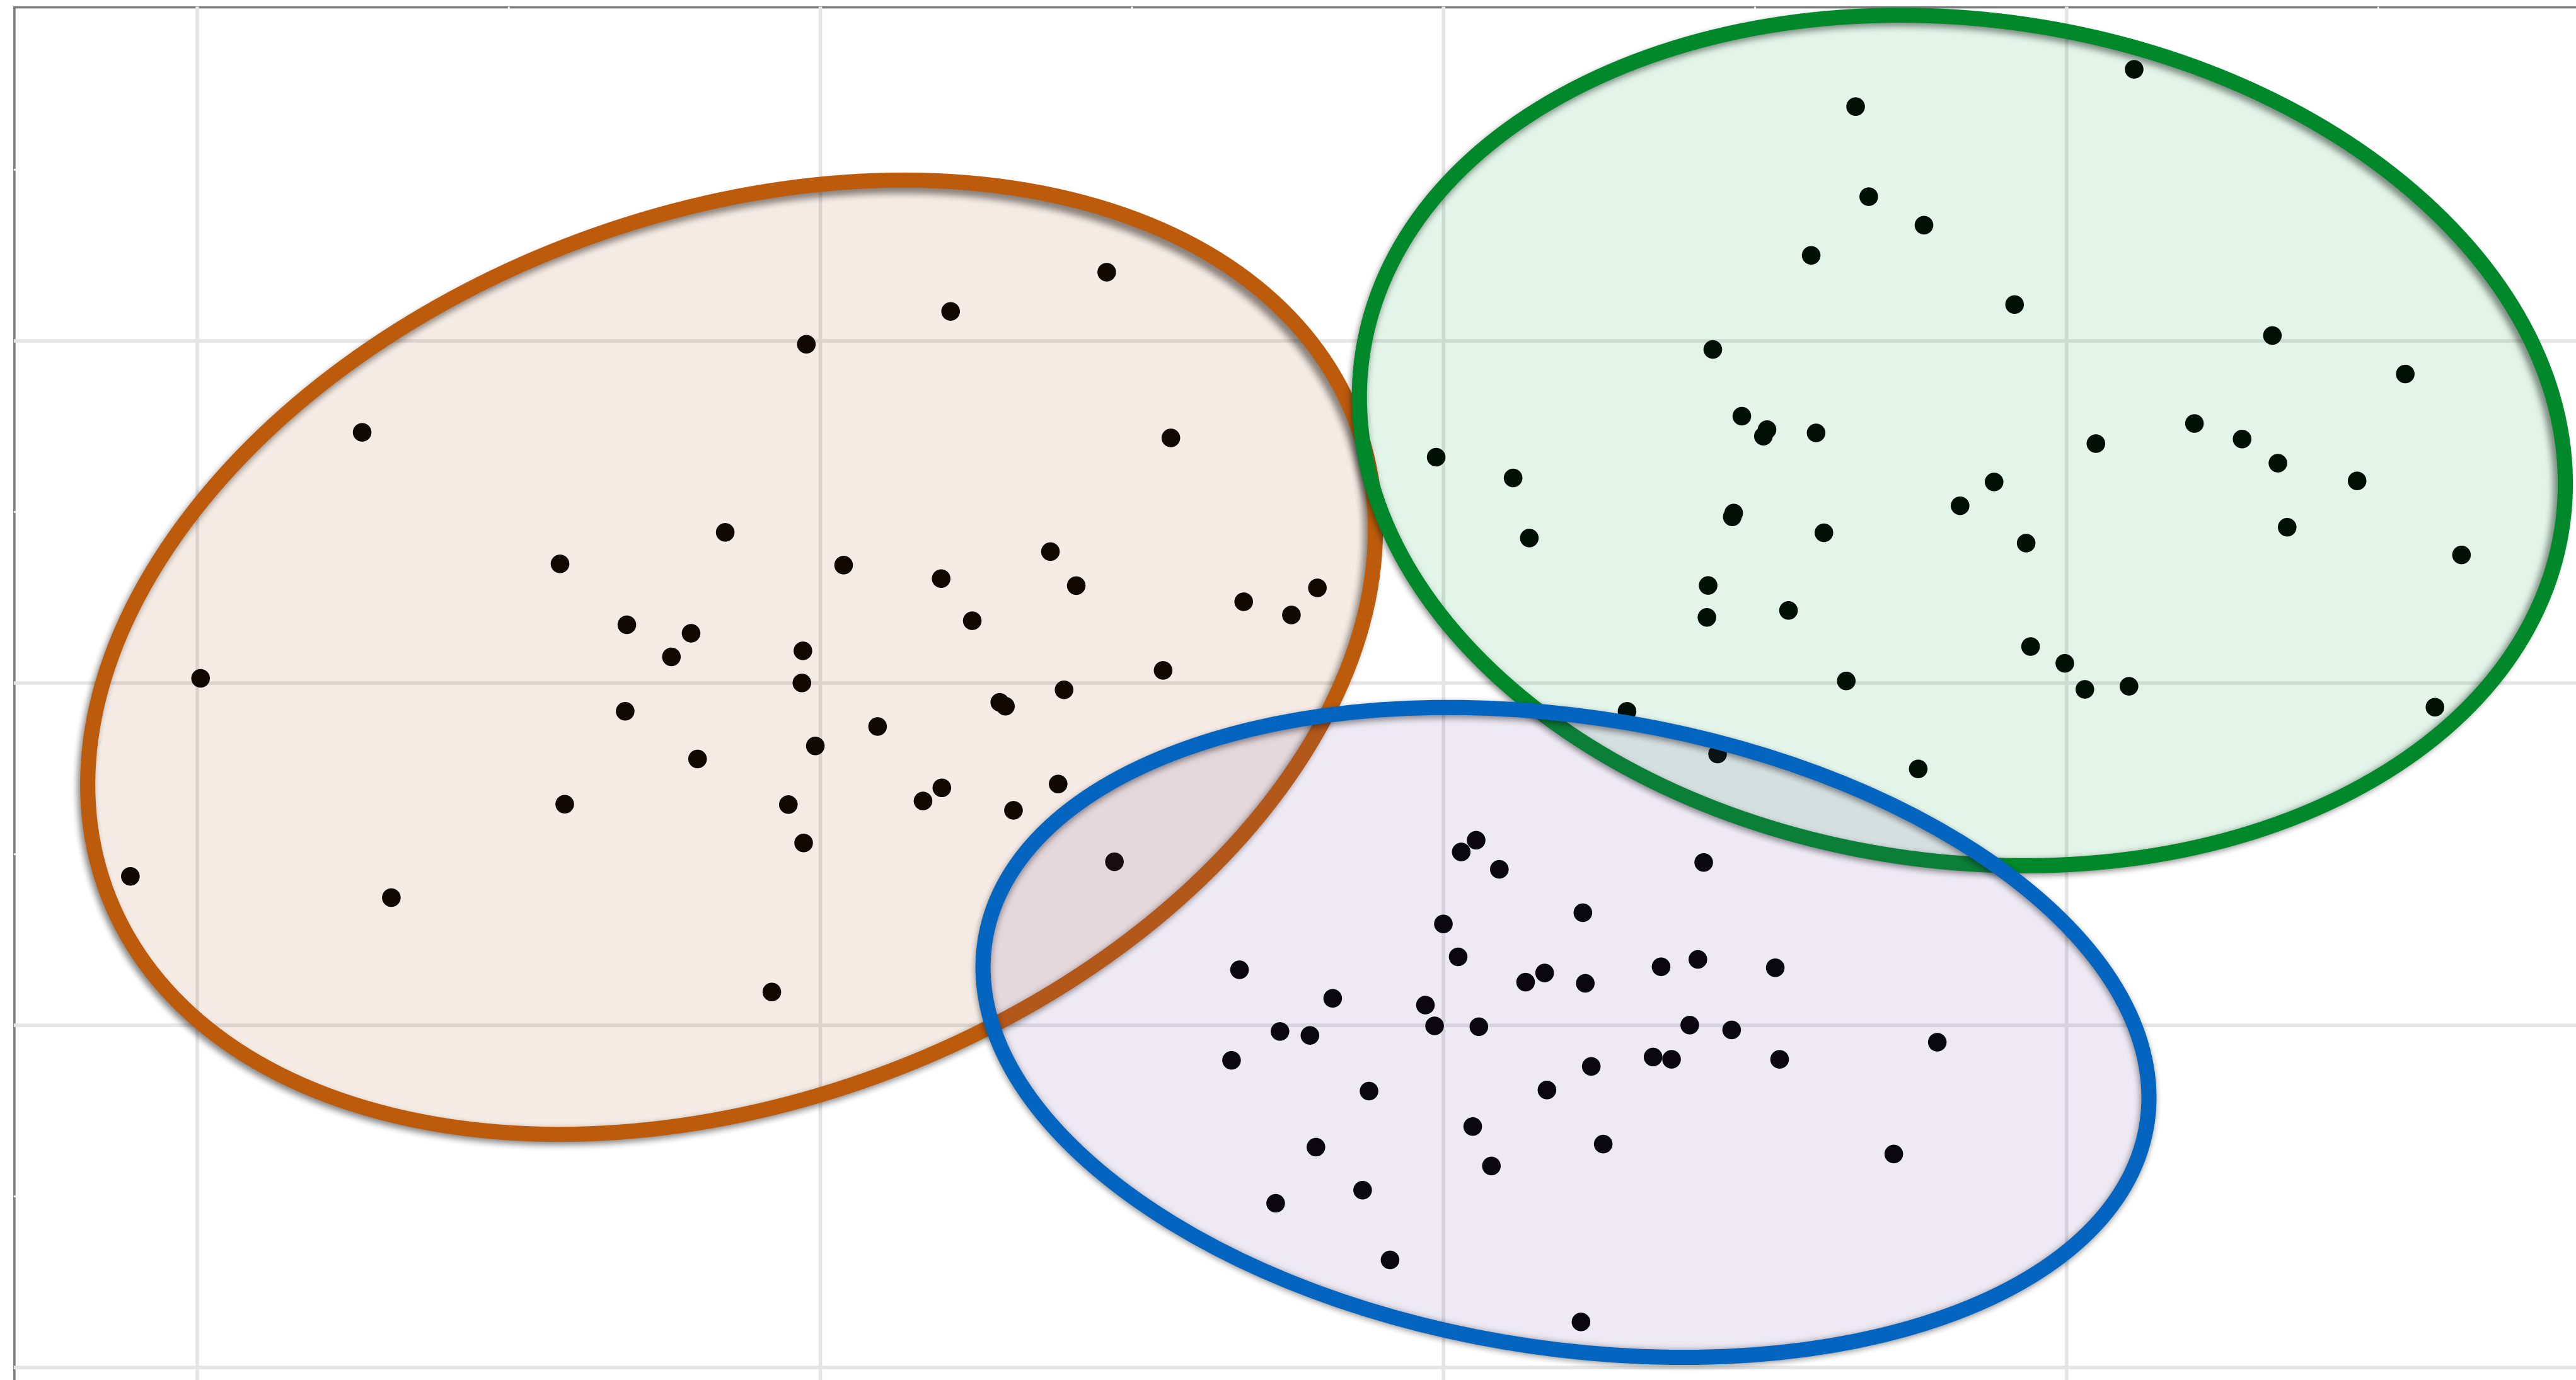
	Malware events	Phishing	Open Tickets	
Dept1	6	6	3	
Dept2	1	2	1	
Dept3	8	1	9	

Stop


Clustering...



Clustering...



Clustering...




K-Means


Before starting, pick the number of clusters, K

1. Pick K random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting


Step 1: Pick 3 random centroids within data range




Step 2: Assign each data point to the nearest centroid (1)




Step 3: Move centroid to center of assigned points (1)




Step 2: Assign each data point to the nearest centroid (2)




Step 3: Move centroid to center of assigned points (2)




Step 2: Assign each data point to the nearest centroid (3)




Step 3: Move centroid to center of assigned points (3)




Step 2: Assign each data point to the nearest centroid (4)




Step 3: Move centroid to center of assigned points (4)




Step 2: Assign each data point to the nearest centroid (S)



Step 3: Move centroid to center of assigned points (5)



Step 3: Move centroid to center of assigned points (5)



Stop

K-Means: Got a problem with it?

Before starting, pick the number of clusters, K


1. Pick K random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting

K-Means: Got a problem with it?


Before starting, pick the number of clusters, K *Subjective*

1. Pick K random centroids within data range *Not Repeatable*
2. Assign each data point to the nearest centroid *Sensitive to Scale*
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting


How many clusters?




Random Start...




Random Start...



Random Start...



K-Means




“...it’s too easy to throw k-means on your data, and nevertheless get a result out (that is pretty much random, but you won't notice).”


— Anony-Mousse

<http://stats.stackexchange.com/questions/133656/how-to-understand-the-drawbacks-of-k-means>

Which one is correct?





A




B

Which one is correct?



Pain of optimization...Being stuck at sub-optimal local minimum...



Initial guess matters!
Same outcome cannot be guaranteed


Stop

K-Means in practice (Python version)

```
#Import from Scikit-learn
from sklearn.cluster import KMeans

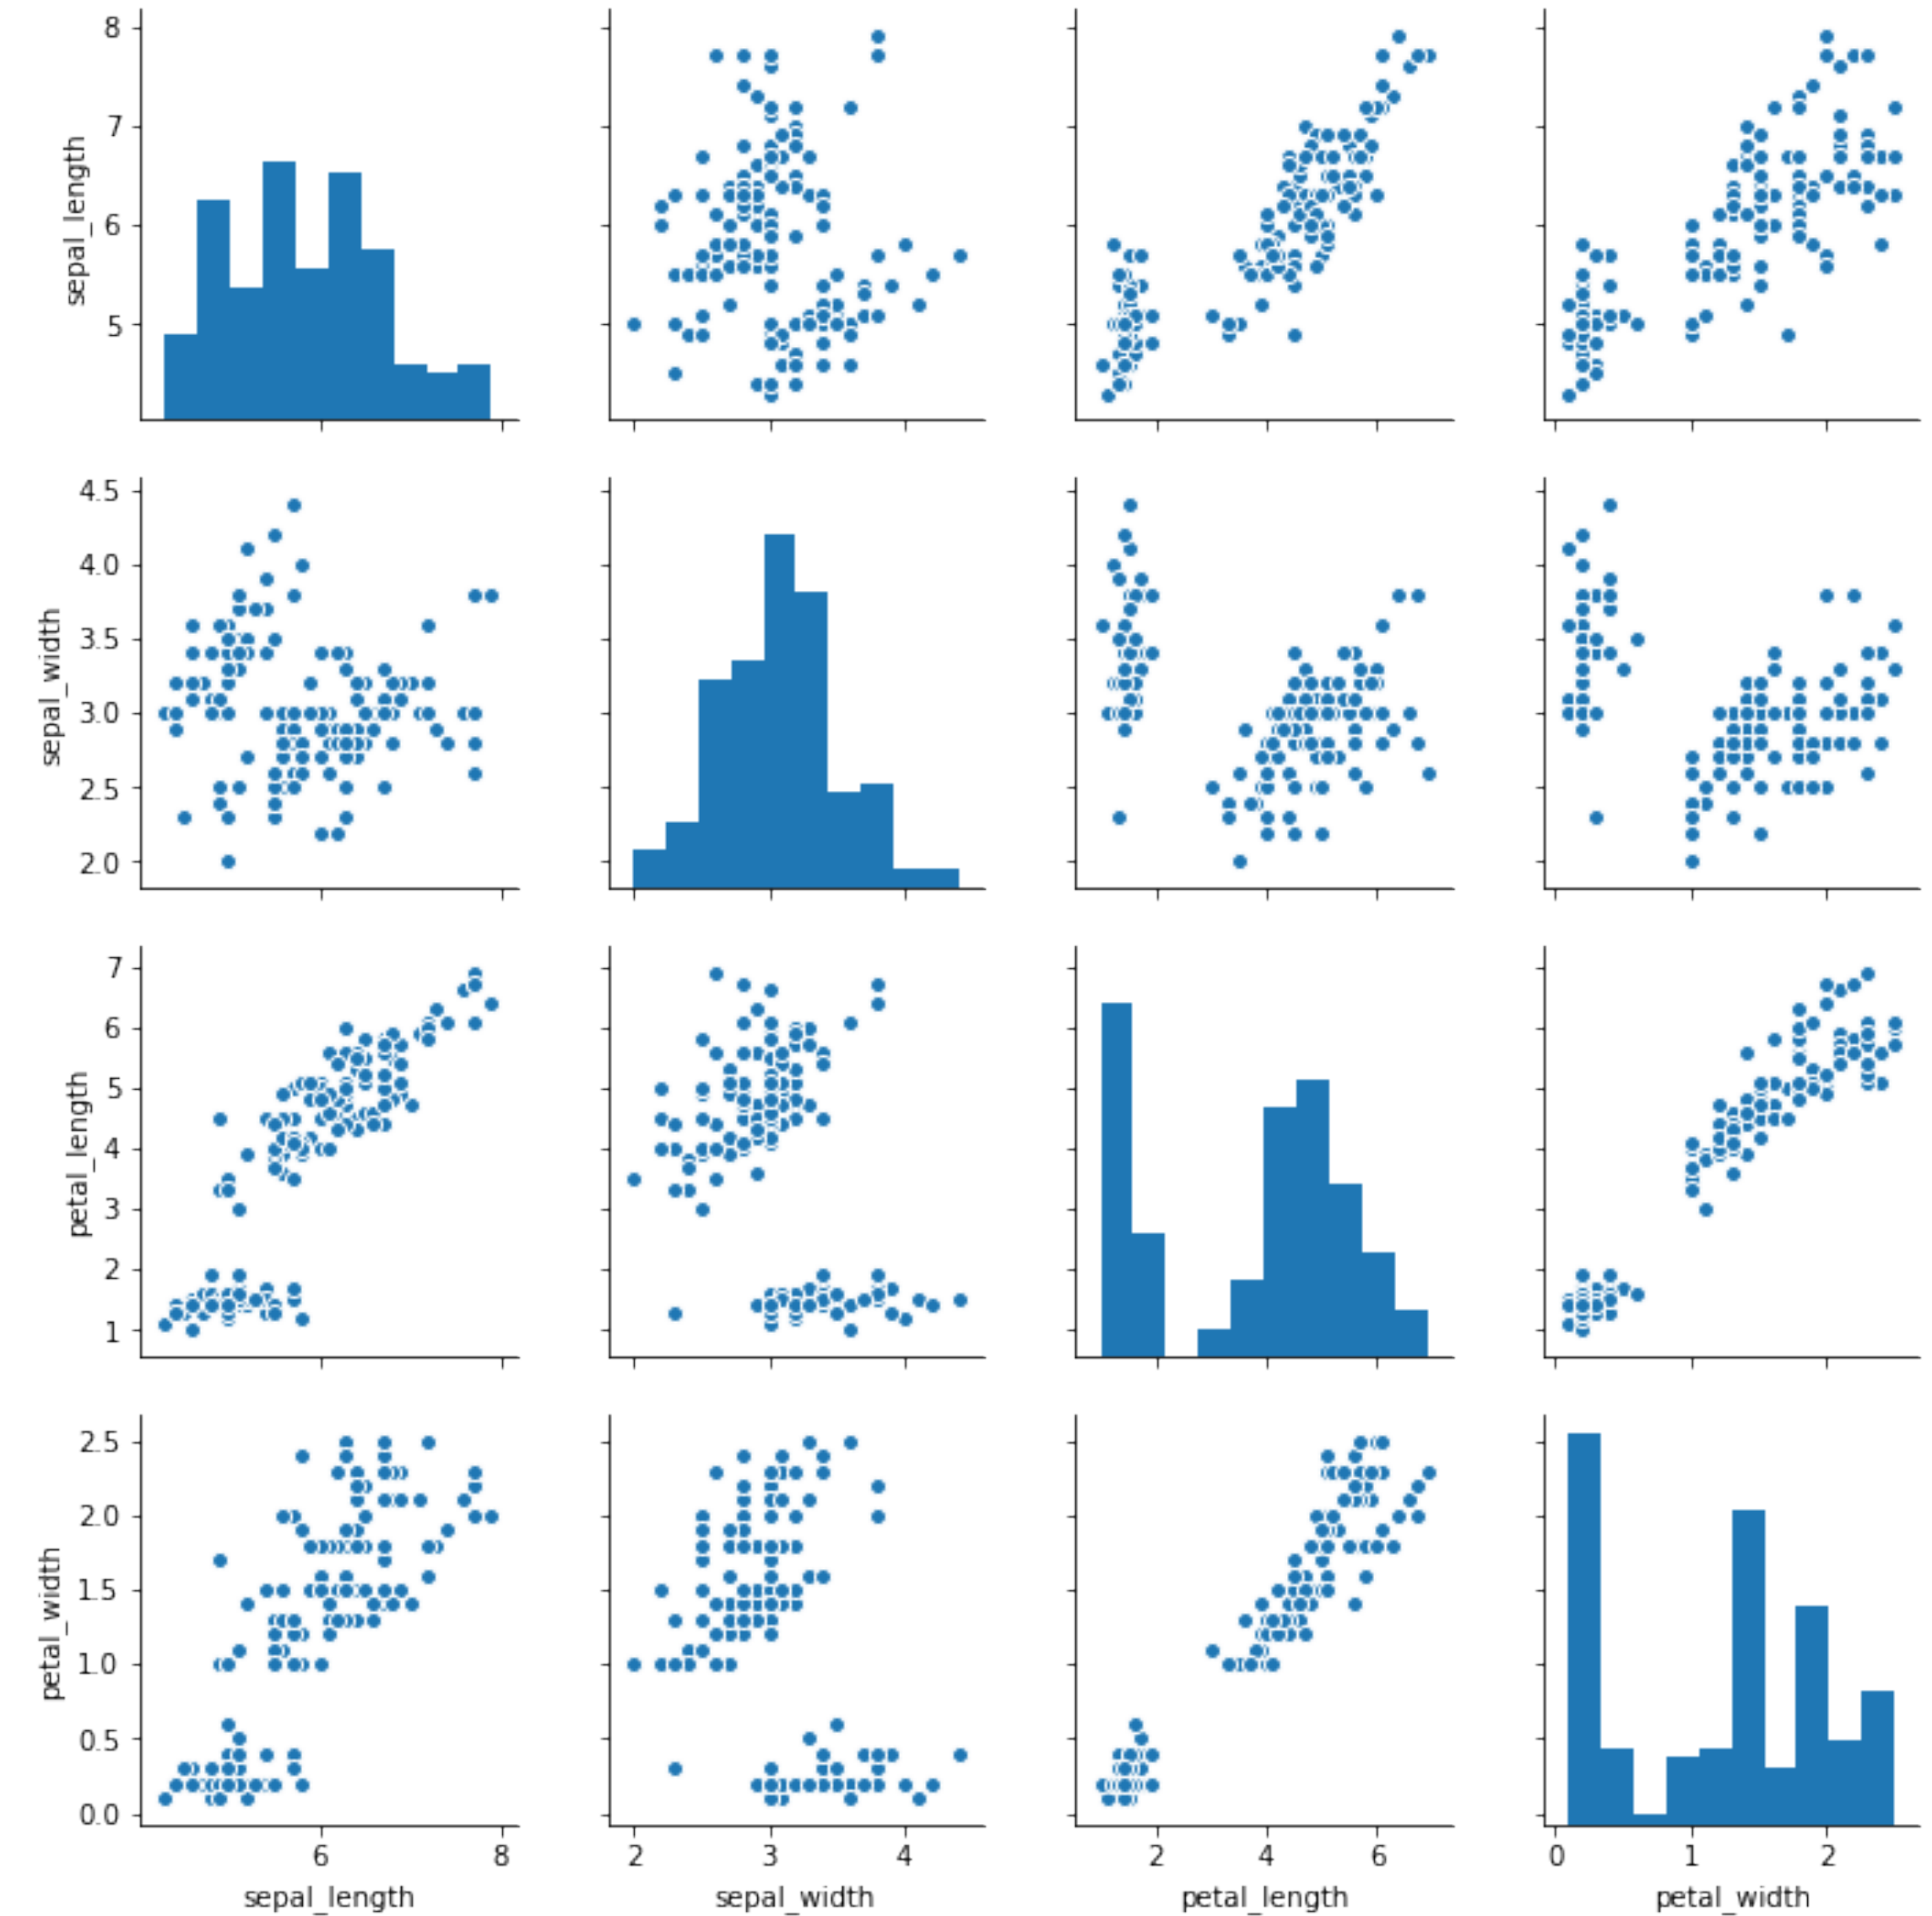
kmeans = KMeans(n_clusters=2)
kmeans.fit(data)

centroids = kmeans.cluster_centers_
labels = kmeans.labels_f
```



The Dataset

```
sns.pairplot(<data>)
```




K-Means Clustering

```
kmeans = KMeans( n_clusters=2 )  
kmeans.fit( <data> )
```

K-Means Clustering

```
sns.pairplot(<data>, x_vars="col_1", y_vars="col_2", hue="kmeans_2", size=6)  
plt.scatter(<cluster_centers>, <col_2>, linewidths=3, marker='x', s=200,  
c='black' )
```



**K-Means is affected by the scale
of every feature.**

Feature Scaling

For k-means clustering, features must be scaled to the same ranges of values to contribute "equally" to the euclidean distance calculation.


Each row is transformed per-column by:

- Subtracting from the element in each row the mean for each feature (column) and then taking this value and
- Dividing by that feature's (column's) standard deviation.

Evaluating your Model

```
WCSS=[ ]  
for i in range(1,20):  
    kmeans=KMeans(n_clusters=I, init='k-means++')  
  
kmeans.fit(final_data[scaled_feature_columns].sample(50000))  
WCSS.append(kmeans.inertia_)
```

Evaluating your Model



Evaluating your Model

```
from yellowbrick.cluster import SilhouetteVisualizer  
model = MiniBatchKMeans(6)  
visualizer = SilhouetteVisualizer(model)
```

```
visualizer.fit(X)  
visualizer.poof()
```

